

Analysis of PennSound speech-to-text

includes additional analysis

1 using DataFrames, CSV, Statistics, Dates, Plots, Distributions,RollingFunctions

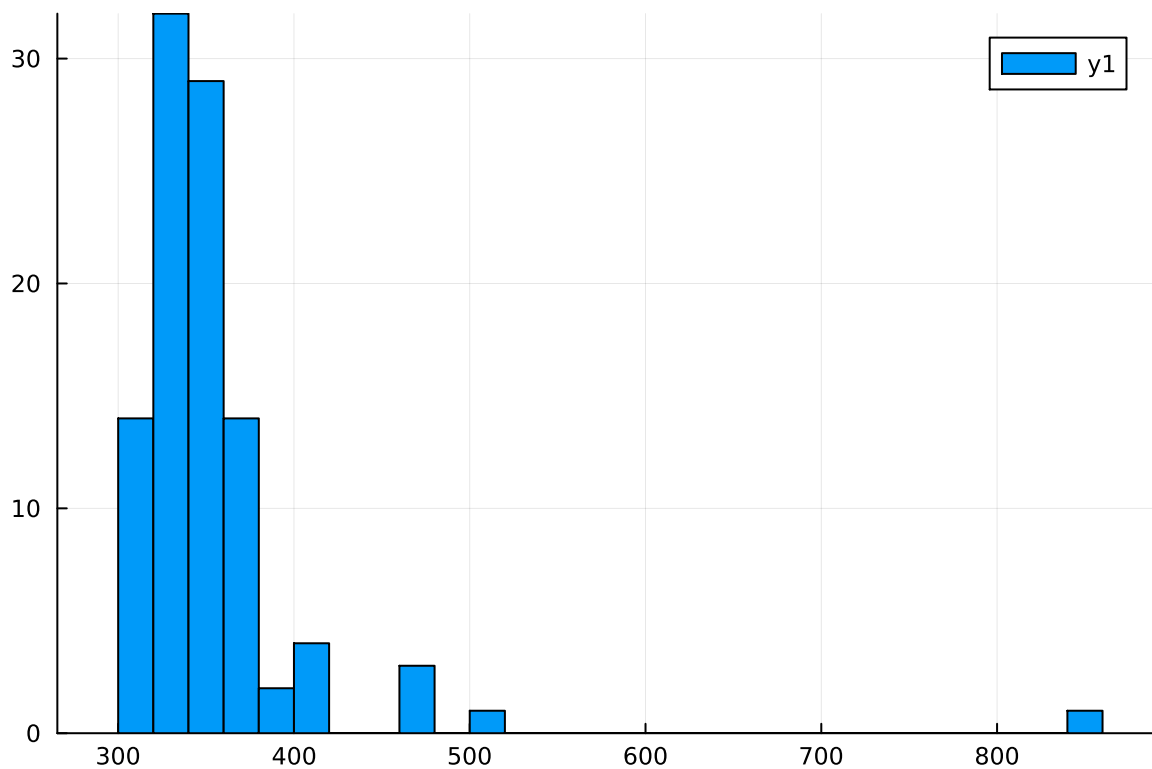
Sample Characteristics

The original goal was a random sample of 100 clips, each 5 minutes in duration, which would be 8.3 hours of audio if done exactly that way. The total audio duration is 12.15 hours, which would mean quite a lot of silence if the speech was only 8.3 hours. However, the reference transcripts tend to indicate more than 5 minutes of speech, due to speech undetected by SAD, overlapping speech, or possibly human error (segments padded with silence). Summing the segment lengths from the human transcripts gives a total of 9.84 hours of speech, with a mean of 354 seconds and a standard deviation of 61.5 seconds. The histogram below shows the amount of speech per file in seconds.

speech =

	file	speech
1	"Andrews-Bruce-and-Charles-North_Compl	341.623
2	"Antin-David_Complete_Seminar_Universi	372.994
3	"Ashbery-John_01_Complete-Reading_WBAI	343.004
4	"Ashbery-John_Complete-Reading_Contemp	348.673
5	"Ashbery-John_Complete-Recording_Attit	300.015
6	"Ashbery-John_Complete-Recording_Honor	361.438
7	"Ashbery-John_Complete-Recording_Pione	315.977
8	"Ashbery-John_Complete-Recording_St-Ma	354.344
9	"Ashbery-John_Complete-Recording_The-S	336.321
10	"Ashbery-John_Complete-Recording_WBAI-	363.231
more		
100	"Yau-John_02_Complete-Reading_SUNY-Buf	337.962

```
1 speech = CSV.read("speech.tsv", DataFrame, delim="\t")
2
```



```
1 histogram(speech[:, :speech])
```

354.16752999999994

```
1 mean(speech.speech)
```

61.48031766484953

```
1 std(speech.speech)
```

file	speech
1	"Ginsberg-Allen_Complete-Reading_WCW-L 858.552

```
1 speech[speech.speech .> 800, :]
```

9.837986944444443

```
1 sum(speech.speech) / 3600
```

total duration of audio clips is 12.15 hours

durations =

	file	duration
1	"Andrews-Bruce-and-Charles-North_Compl	585.15
2	"Antin-David_Complete_Seminar_Universi	462.47
3	"Ashbery-John_01_Complete-Reading_WBAI	432.82
4	"Ashbery-John_Complete-Reading_Contemp	446.79
5	"Ashbery-John_Complete-Recording_Attit	390.93
6	"Ashbery-John_Complete-Recording_Honor	458.35
7	"Ashbery-John_Complete-Recording_Pione	393.12
8	"Ashbery-John_Complete-Recording_St-Ma	466.21
9	"Ashbery-John_Complete-Recording_The-S	440.09
10	"Ashbery-John_Complete-Recording_WBAI-	458.39
more		
100	"Yau-John_02_Complete-Reading_SUNY-Buf	379.11

```
1 durations = CSV.read("durations.tsv", DataFrame, delim="\t")
2
```

12.152402777777779

```
1 sum(durations.duration) / 3600
```

Word Error Rates

here we display individual WERs in various ways

	vs	azure	google	ibm	nemo	rev	whisper	whispercpp	nsp	snr
1	0	23.1	24.7	33.4	18.8	19.0	19.1	21.6	1	33.63
2	7	18.1	16.4	21.2	20.3	15.0	19.4	21.2	8	6.94
3		4.3	4.2	5.8	3.4	3.7	2.8	3.5	1	20.8
4		4.6	4.2	6.6	4.4	3.2	3.8	4.6	1	20.02
5		4.8	5.5	6.2	4.9	3.2	4.7	5.3	2	7.43
6	5	14.9	15.6	17.4	18.2	15.6	18.1	17.2	2	21.63
7		6.3	4.8	9.1	5.7	3.2	3.9	4.7	1	14.83
8		2.4	4.1	3.6	2.8	2.4	2.0	6.4	1	19.02
9		4.8	5.2	11.4	3.5	4.2	3.0	3.6	1	-0.31
10		4.5	4.6	8.6	4.7	4.2	3.5	5.8	1	12.39
more										
100		6.0	7.4	10.7	4.4	4.5	3.6	4.9	1	18.16

```
1 wers = CSV.read("wer.tsv", DataFrame, delim="\t")
2
```

substitution error rates only

	file	aws	azure	google	ibm	nemo	rev	whisper	whispercpp
66	erbach1"	5.1	4.2	4.5	7.2	3.1	4.8	2.6	3.5
67	erbach2"	3.1	3.4	4.3	5.7	2.7	4.6	3.0	2.9
68	'	4.7	4.7	5.4	7.7	4.5	4.8	3.9	4.7
69	ove"	4.4	4.0	5.1	5.7	3.4	3.6	3.3	3.8
70	e"	2.9	3.7	4.0	6.9	2.9	3.2	2.6	2.9
71	ey"	9.1	11.8	9.8	18.4	7.4	8.4	9.1	8.6
72	i"	11.5	13.1	12.1	17.7	10.4	10.6	10.3	10.4
73	ytalks1"	5.1	4.4	4.6	7.6	3.6	5.1	4.3	4.8
74	ytalks10"	4.7	5.9	4.2	6.7	3.0	2.9	3.5	3.3
75	ytalks3"	4.5	4.5	4.7	6.6	3.1	4.0	2.0	2.7
76	ytalks5"	2.7	2.0	1.9	3.1	1.5	2.6	1.6	1.9
77	ino"	2.7	2.0	2.8	4.5	1.2	3.3	2.4	2.3
78	alk"	3.7	4.1	4.5	6.5	3.8	4.1	3.2	2.8
79	th"	9.4	9.3	11.3	15.2	11.3	10.4	9.1	9.0
80	ack"	10.9	9.6	9.3	15.0	8.5	6.7	7.0	7.9
81	lkoff"	4.5	3.4	4.7	7.8	2.5	3.2	2.5	2.0
82	ards"	3.1	3.1	2.2	7.4	0.6	1.9	1.5	2.3
83	son1"	1.7	1.8	2.7	3.3	1.1	1.3	0.9	1.7

```
1 wers_s = CSV.read("wer_s.tsv", DataFrame, delim="\t")
2
```

deletion error rates only

	file	aws	azure	google	ibm	nemo	rev	whisper	whis
65	"kyger"	10.6	11.3	10.1	12.8	9.3	5.6	7.3	6.7
66	"lauterbach1"	1.2	2.2	3.2	2.4	3.9	2.1	2.2	3.3
67	"lauterbach2"	1.4	2.1	1.6	2.1	5.2	1.2	3.7	7.0
68	"levy"	0.7	0.8	1.8	1.2	2.0	0.9	2.2	1.3
69	"mirakove"	1.0	0.5	1.2	1.3	2.0	0.9	1.1	1.3
70	"moore"	0.0	0.1	1.4	0.5	0.5	0.9	0.0	0.6
71	"moxley"	5.0	3.6	4.5	3.5	4.1	3.8	1.9	3.1
72	"oppen"	3.2	3.8	4.8	5.7	3.7	3.9	1.7	4.0
73	"phillytalks1"	7.8	9.5	9.5	13.3	13.6	6.0	8.6	10.7
74	"phillytalks10"	0.9	1.3	1.1	1.8	1.4	1.0	1.5	1.1
75	"phillytalks3"	1.1	2.7	2.0	2.2	6.4	1.9	5.7	5.9
76	"phillytalks5"	8.1	11.4	8.7	10.4	13.3	5.1	10.2	7.4
77	"piombino"	1.1	0.6	1.5	0.9	1.7	0.6	1.4	2.0
78	"poemtalk"	5.7	4.9	3.9	5.1	9.5	3.2	5.5	5.5
79	"raworth"	0.8	1.3	4.0	2.3	1.4	1.1	1.1	1.6
80	"retalack"	8.0	10.3	10.8	10.9	15.0	10.4	13.9	14.0
81	"reznikoff"	1.2	2.4	2.5	3.1	3.1	3.0	2.1	2.5

```
1 wers_d = CSV.read("wer_d.tsv", DataFrame, delim="\t")
2
```

insertion error rates only

	file	aws	azure	google	ibm	nemo	rev	whisper	whis
83	"robinson1"	0.6	0.6	0.4	0.3	0.4	0.1	0.7	0.2
84	"robinson2"	1.0	0.9	1.0	0.7	0.8	0.5	0.6	0.8
85	"robinson3"	1.2	0.9	0.7	0.8	1.2	0.4	0.5	0.7
86	"rothenberg"	1.1	0.5	1.2	1.2	1.1	0.9	1.1	7.5
87	"scalapino1"	0.8	0.7	0.5	0.9	0.4	0.7	0.5	0.7
88	"scalapino2"	1.4	0.9	1.4	1.2	1.3	0.9	1.8	3.2
89	"sherlock"	1.1	1.4	0.3	1.2	1.0	0.6	0.9	0.8
90	"silliman1"	0.5	0.6	0.7	0.8	0.3	1.0	2.1	1.3
91	"silliman2"	0.8	0.8	0.8	0.8	1.1	1.3	1.4	0.9
92	"silliman3"	0.7	0.7	0.5	1.0	0.5	0.2	0.7	0.8
93	"smith"	2.0	1.2	1.6	1.2	1.2	1.2	0.6	0.6
94	"spahr"	1.5	1.4	1.1	1.2	0.8	0.8	0.9	0.8
95	"sze"	1.3	0.8	1.0	0.8	0.7	0.7	0.8	0.4
96	"templeton"	1.5	1.6	1.0	0.8	2.4	3.5	3.3	3.9
97	"torres"	6.4	5.0	5.3	5.6	3.9	6.7	3.7	12.2
98	"towle"	0.6	0.5	0.7	1.2	0.1	0.2	0.4	0.1
99	"wisher"	0.6	1.2	0.8	1.2	0.8	0.3	0.4	0.3

```
1 wers_i = CSV.read("wer_i.tsv", DataFrame, delim="\t")
2
```

sort by rev, the best performer

wers_sorted_by_rev =

	file	aws	azure	google	ibm	nemo	rev	whisper	whisperc
1	"bromige2"	2.4	3.0	2.0	4.2	2.4	1.1	2.0	1.9
2	"duplessis1"	3.0	2.9	2.9	4.4	3.9	1.8	1.9	5.4
3	"duplessis2"	2.8	2.6	3.6	6.7	2.9	2.0	1.8	1.6
4	"gladman"	2.2	2.4	3.2	4.1	4.2	2.2	1.9	10.9
5	"jarnot"	4.3	4.0	4.5	9.3	5.5	2.3	3.1	4.3
6	"ashbery6"	2.0	2.4	4.1	3.6	2.8	2.4	2.0	6.4
7	"robinson1"	3.0	2.8	5.0	4.2	4.1	2.5	2.6	3.6
8	"richards"	3.9	3.9	2.8	8.2	2.8	2.8	3.4	7.1
9	"robinson3"	3.6	3.2	3.9	5.2	4.2	2.9	2.9	5.3
10	"ashbery2"	3.6	4.6	4.2	6.6	4.4	3.2	3.8	4.6
more									
100	"ginsberg"	37.2	36.3	36.2	44.0	34.4	33.3	32.4	39.4

```
1 wers_sorted_by_rev = sort(wers, [:rev])
```

add mean WER to table

wers_with_mean =

	file	aws	azure	google	ibm	nemo	rev	whisper	whispercpp
1	"andrews"	29.0	23.1	24.7	33.4	18.8	19.0	19.1	21.6
2	"antin"	15.7	18.1	16.4	21.2	20.3	15.0	19.4	21.2
3	"ashbery1"	3.5	4.3	4.2	5.8	3.4	3.7	2.8	3.5
4	"ashbery2"	3.6	4.6	4.2	6.6	4.4	3.2	3.8	4.6
5	"ashbery3"	3.9	4.8	5.5	6.2	4.9	3.2	4.7	5.3
6	"ashbery4"	14.6	14.9	15.6	17.4	18.2	15.6	18.1	17.2
7	"ashbery5"	5.4	6.3	4.8	9.1	5.7	3.2	3.9	4.7
8	"ashbery6"	2.0	2.4	4.1	3.6	2.8	2.4	2.0	6.4
9	"ashbery7"	4.1	4.8	5.2	11.4	3.5	4.2	3.0	3.6
10	"ashbery8"	4.5	4.5	4.6	8.6	4.7	4.2	3.5	5.8
more									
100	"yau"	5.5	6.0	7.4	10.7	4.4	4.5	3.6	4.9

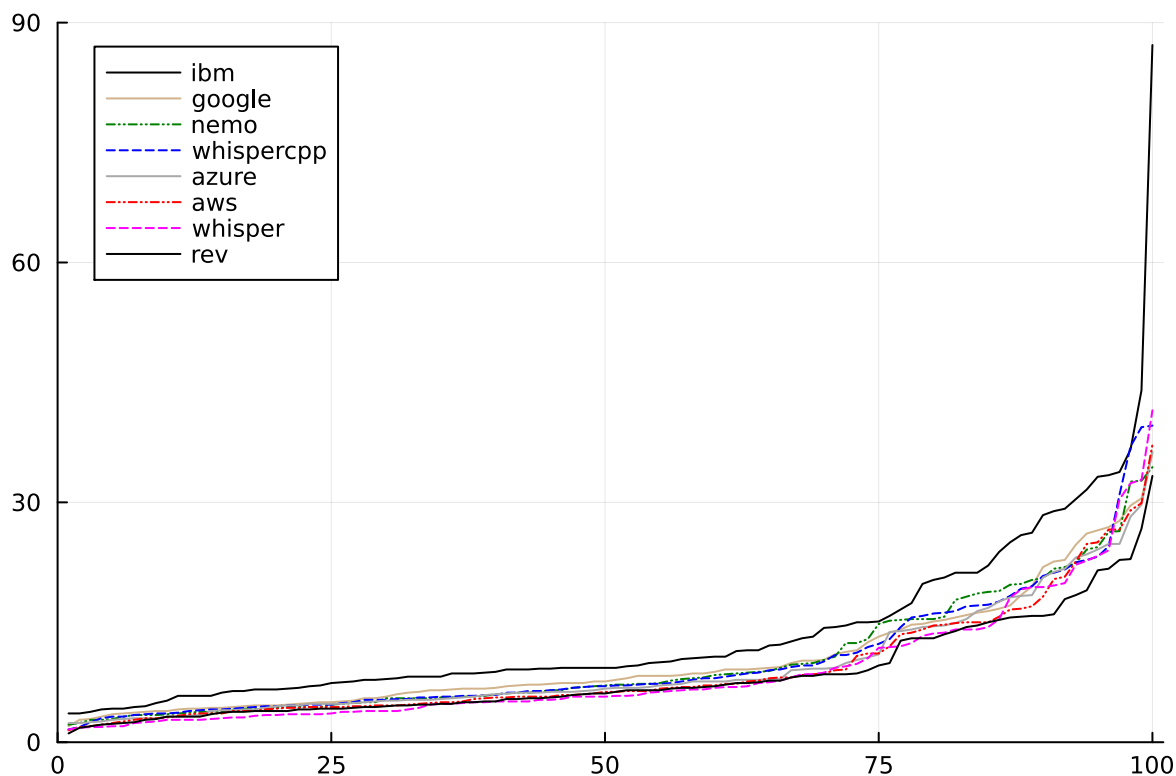
```
1 wers_with_mean = transform(wers, AsTable([:aws, :azure, :google, :ibm, :nemo, :rev, :whisper, :whispercpp]) => ByRow(mean) => :mean)
```

sort by mean WER

	file	aws	azure	google	ibm	nemo	rev	whisper	whis
25	"ashbery7"	4.1	4.8	5.2	11.4	3.5	4.2	3.0	3.6
26	"fiedler"	5.0	5.3	4.3	8.2	6.2	3.8	3.4	4.0
27	"ashbery8"	4.5	4.5	4.6	8.6	4.7	4.2	3.5	5.8
28	"howe1"	5.9	4.8	4.3	10.1	4.8	3.9	3.9	3.4
29	"coolidge"	4.2	5.4	6.6	6.6	5.6	5.1	3.5	5.3
30	"ashbery5"	5.4	6.3	4.8	9.1	5.7	3.2	3.9	4.7
31	"drucker2"	4.6	5.3	7.0	7.0	3.6	6.1	3.8	5.7
32	"hawkins"	3.3	3.0	4.9	3.6	3.8	4.1	3.1	17.6
33	"wisher"	4.7	5.3	6.0	7.8	6.2	5.4	3.9	4.5
34	"bellamy"	4.5	4.6	8.6	5.8	4.4	5.5	5.0	6.4
35	"sze"	6.9	5.1	6.8	8.7	5.5	5.5	3.9	4.4
36	"yau"	5.5	6.0	7.4	10.7	4.4	4.5	3.6	4.9
37	"mirakove"	5.7	5.2	7.4	7.6	6.3	5.1	4.8	5.4
38	"davies"	5.6	5.0	5.0	8.8	5.5	4.7	4.9	8.7
39	"silliman1"	4.6	5.9	5.1	9.3	7.0	4.6	6.5	6.2
40	"garrison"	6.0	6.2	6.5	7.9	5.5	6.8	4.8	5.9
41	"bromige4"	4.9	6.9	6.5	9.2	5.8	5.0	5.7	5.6

```
1 wers_sorted_by_mean = sort(wers_with_mean, :mean)
```

WERs sorted separately

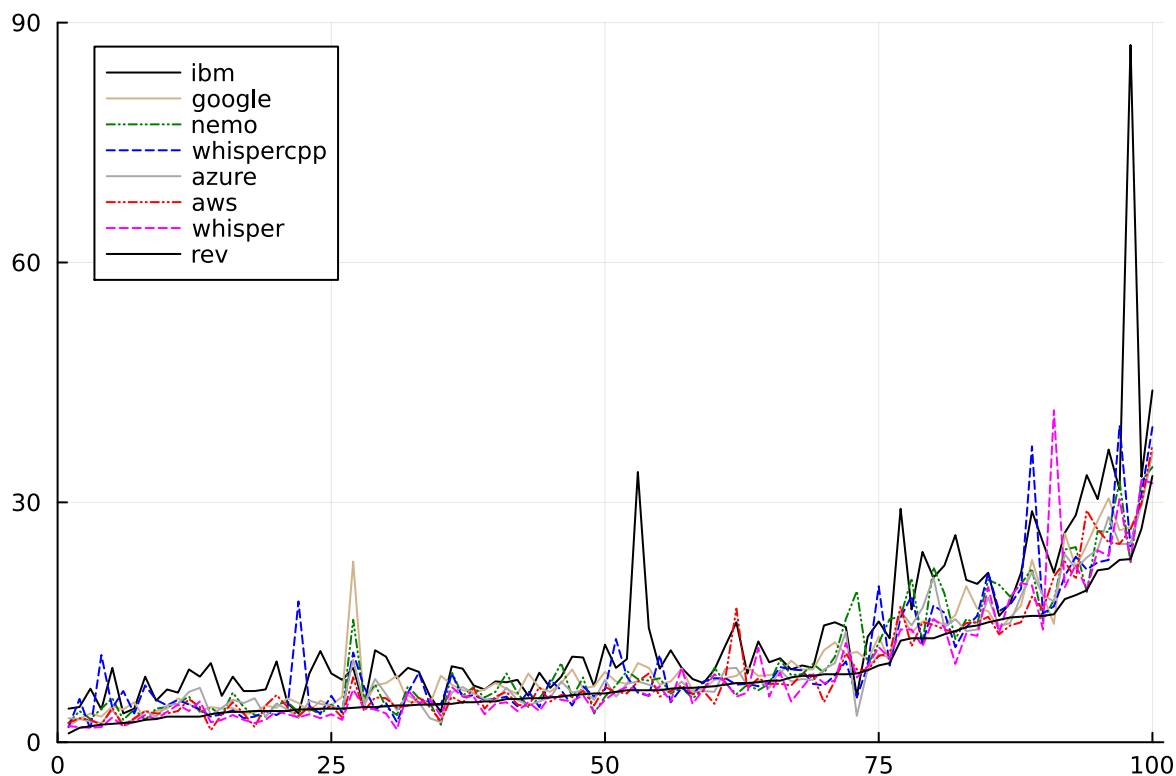


```

1 begin
2   plot(sort(wers.ibm),label="ibm", linestyle=:solid, color=:black)
3   plot!(sort(wers.google),label="google", linestyle=:solid, color=:tan)
4   plot!(sort(wers.nemo),label="nemo", linestyle=:dashdotdot, color=:green)
5   plot!(sort(wers.whispercpp),label="whispercpp", linestyle=:dash, color=:blue)
6   plot!(sort(wers.azure),label="azure", linestyle=:solid, color=:darkgray)
7   plot!(sort(wers.aws),label="aws",xlim=(0,101),ylim=(0,90),
8     linestyle=:dashdotdot,color=:red)
9   plot!(sort(wers.whisper),label="whisper", linestyle=:dash, color=:magenta)
10  plot!(sort(wers.rev),label="rev", linestyle=:solid,color=:black)
10 end

```

WERs sorted by rev (best system)

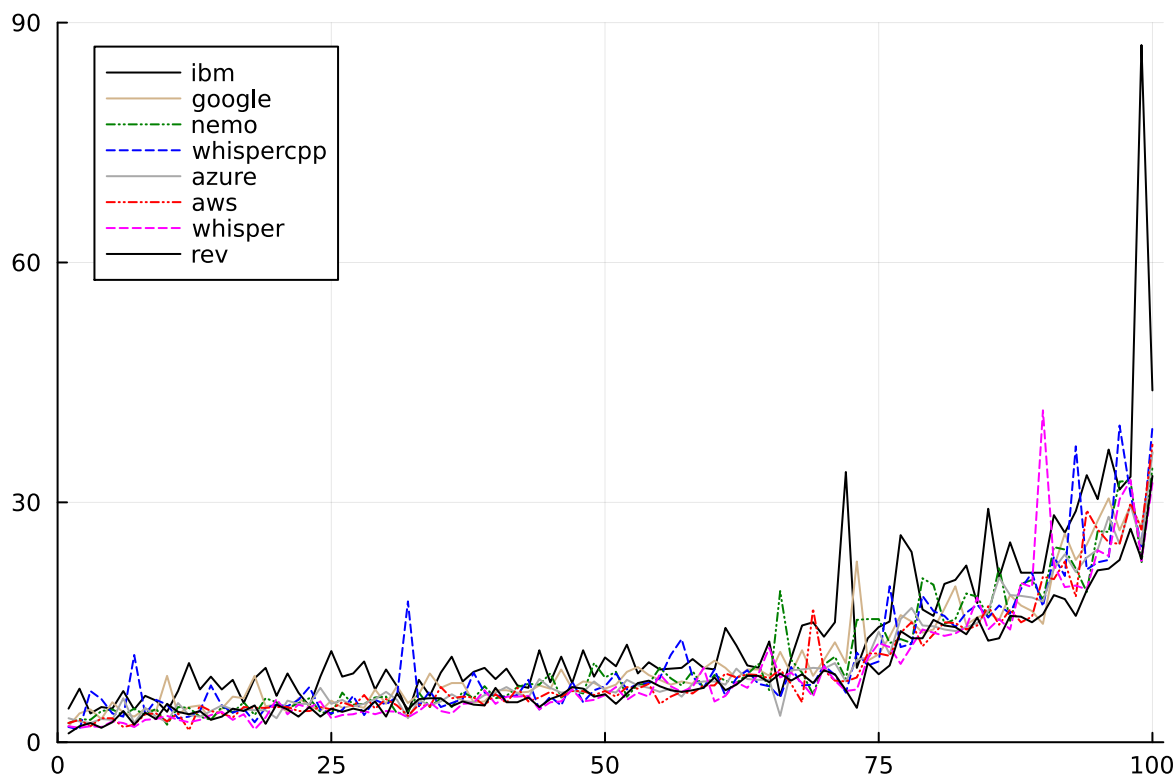


```

1 begin
2   w1 = wers_sorted_by_rev
3   plot(w1.ibm,label="ibm", linestyle=:solid, color=:black)
4   plot!(w1.google,label="google", linestyle=:solid, color=:tan)
5   plot!(w1.nemo,label="nemo", linestyle=:dashdotdot, color=:green)
6   plot!(w1.whispercpp,label="whispercpp", linestyle=:dash, color=:blue)
7   plot!(w1.azure,label="azure", linestyle=:solid, color=:darkgray)
8   plot!(w1.aws,label="aws",xlim=(0,101),ylim=(0,90),
9         linestyle=:dashdotdot,color=:red)
10  plot!(w1.whisper,label="whisper", linestyle=:dash, color=:magenta)
11  plot!(w1.rev,label="rev", linestyle=:solid,color=:black)
12 end

```

WERs sorted by mean WER



```

1 begin
2   w2 = wers_sorted_by_mean
3   plot(w2.ibm,label="ibm", linestyle=:solid, color=:black)
4   plot!(w2.google,label="google", linestyle=:solid, color=:tan)
5   plot!(w2.nemo,label="nemo", linestyle=:dashdotdot, color=:green)
6   plot!(w2.whispercpp,label="whispercpp", linestyle=:dash, color=:blue)
7   plot!(w2.azure,label="azure", linestyle=:solid, color=:darkgray)
8   plot!(w2.aws,label="aws",xlim=(0,101),ylim=(0,90),
9         linestyle=:dashdotdot,color=:red)
10  plot!(w2.whisper,label="whisper", linestyle=:dash, color=:magenta)
11  plot!(w2.rev,label="rev", linestyle=:solid,color=:black)
12 end

```

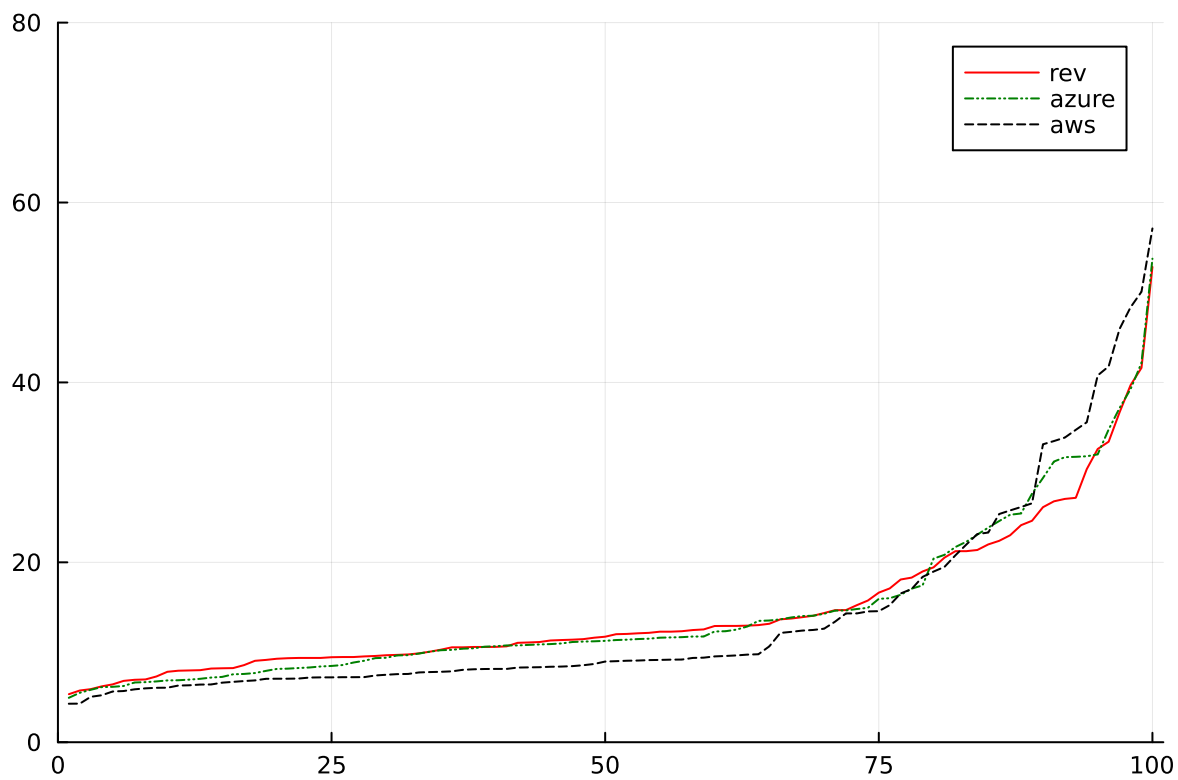
Diarization Error Rates

der =

	file	aws	azure	ibm	rev
1	"Andrews-Bruce-and-Charles-North_Compl	16.54	13.53	62.37	23.01
2	"Antin-David_Complete_Seminar_Universi	20.78	25.43	65.45	27.17
3	"Ashbery-John_01_Complete-Reading_WBAI	9.59	10.62	15.08	10.59
4	"Ashbery-John_Complete-Reading_Contemp	8.15	11.36	47.78	12.47
5	"Ashbery-John_Complete-Recording_Attit	8.05	10.87	18.4	11.73
6	"Ashbery-John_Complete-Recording_Honor	21.98	31.79	67.19	24.62
7	"Ashbery-John_Complete-Recording_Pione	8.41	8.18	20.46	6.82
8	"Ashbery-John_Complete-Recording_St-Ma	9.08	11.47	53.98	9.37
9	"Ashbery-John_Complete-Recording_The-S	8.15	8.48	59.48	9.37
10	"Ashbery-John_Complete-Recording_WBAI-	8.97	11.27	23.86	12.54
more					
100	"Yau-John_02_Complete-Reading_SUNY-Buf	7.06	7.69	8.13	9.47

```
1 der = CSV.read("der.tsv", DataFrame, delim="\t")
2
```

DERs sorted separately



```
1 begin
2   plot(sort(der.rev),label="rev",linestyle=:solid,color=:red,xlim=(0,101),ylim=
3     (0,80))
4   plot!(sort(der.azure),label="azure", linestyle=:dashdotdot, color=:green)
5   plot!(sort(der.aws),label="aws", linestyle=:dash, color=:black)
6 end
```

What about the outliers?

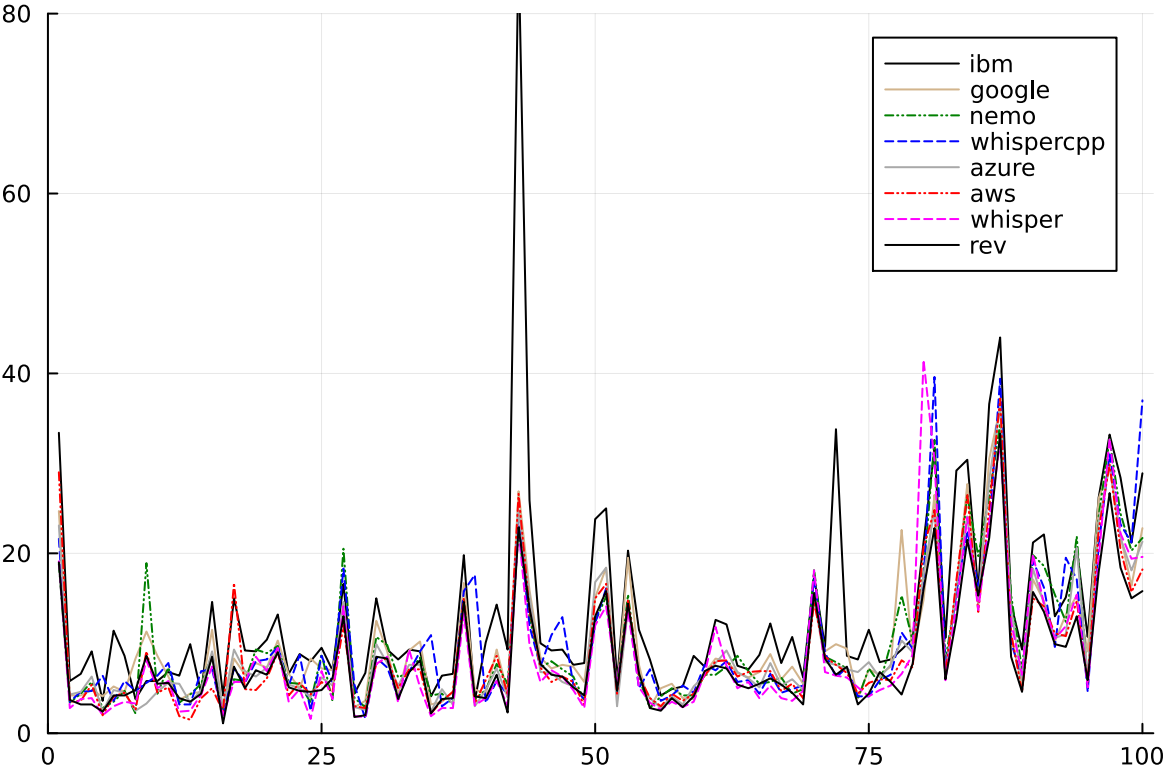
WERs sorted by SNR provided by IBM

wers_sorted_by_snr =

	file	aws	azure	google	ibm	nemo	rev	whisper	whispercpp
1	"andrews"	29.0	23.1	24.7	33.4	18.8	19.0	19.1	21.6
2	"ashbery1"	3.5	4.3	4.2	5.8	3.4	3.7	2.8	3.5
3	"ashbery2"	3.6	4.6	4.2	6.6	4.4	3.2	3.8	4.6
4	"ashbery5"	5.4	6.3	4.8	9.1	5.7	3.2	3.9	4.7
5	"ashbery6"	2.0	2.4	4.1	3.6	2.8	2.4	2.0	6.4
6	"ashbery7"	4.1	4.8	5.2	11.4	3.5	4.2	3.0	3.6
7	"ashbery8"	4.5	4.5	4.6	8.6	4.7	4.2	3.5	5.8
8	"auster"	2.6	2.5	8.3	3.8	2.2	4.8	3.3	4.7
9	"beaulieu"	9.1	3.3	11.3	5.8	18.9	8.6	8.1	5.6
10	"bellamy"	4.5	4.6	8.6	5.8	4.4	5.5	5.0	6.4
more									
100	"benson2"	18.2	21.3	22.8	28.9	21.7	15.8	19.6	37.0

```
1 wers_sorted_by_snr = sort(wers, [:nsp])
```

SNR seems to have an overall effect, but not consistently, and doesn't seem to explain the outliers



```
1 begin
2   plot(wers_sorted_by_snr.ibm,label="ibm", linestyle=:solid, color=:black)
3   plot!(wers_sorted_by_snr.google,label="google", linestyle=:solid, color=:tan)
4   plot!(wers_sorted_by_snr.nemo,label="nemo", linestyle=:dashdotdot, color=:green)
5   plot!(wers_sorted_by_snr.whispercpp,label="whispercpp", linestyle=:dash,
6         color=:blue)
7   plot!(wers_sorted_by_snr.azure,label="azure", linestyle=:solid, color=:darkgray)
8   plot!(wers_sorted_by_snr.aws,label="aws",xlim=(0,101),ylim=(0,80),
9         linestyle=:dashdotdot,color=:red)
10  plot!(wers_sorted_by_snr.whisper,label="whisper", linestyle=:dash,
11        color=:magenta)
12  plot!(wers_sorted_by_snr.rev,label="rev", linestyle=:solid,color=:black)
13 end
```

IBM specifically

what happened with ibm?

There's one extreme outlier, *joris*, where most words are missing. This recording has a lot of distortion/feedback.

	file	aws	azure	google	ibm	nemo	rev	whisper	whispercpp	ns
1	"joris"	26.6	24.8	26.9	87.2	22.4	22.9	22.7	24.5	1

```
1 wers[ wers.file .== "joris", :]
```

Mark, Neville, James, You can listen to it [here](#)

Another outlier is *corrigan*

	file	aws	azure	google	ibm	nemo	rev	whisper	whispercpp
1	"corrigan"	7.6	7.6	9.9	33.8	7.9	6.5	6.4	6.2

```
1 wers[wers.file == "corrigan", :]
```

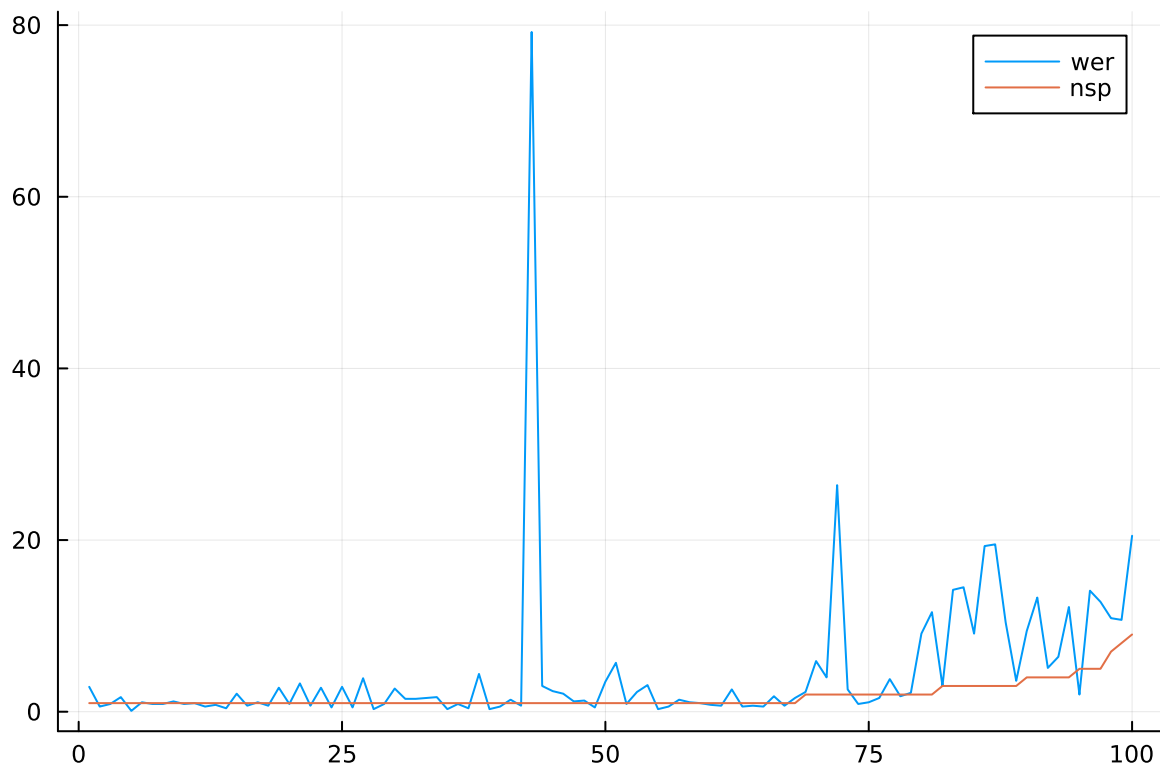
this file is unremarkable, but also has a lot of deletions. E.g., there's a 15s region from 100s to 115s that's missing, which is where a second speaker joins.

	file	aws	azure	google	ibm	nemo	rev	whisper	whispercpp
1	"corrigan"	1.6	1.4	3.0	26.4	3.3	1.3	1.3	1.4

```
1 wers_d[ wers_d.file == "corrigan", :]
```

so IBM seems to have an issue with deletions, maybe due to speakers. sorting by number of speakers shows some effect. The first 68 have a single speaker. *corrigan* has two speakers, so this doesn't seem total explanatory for it's high DER.

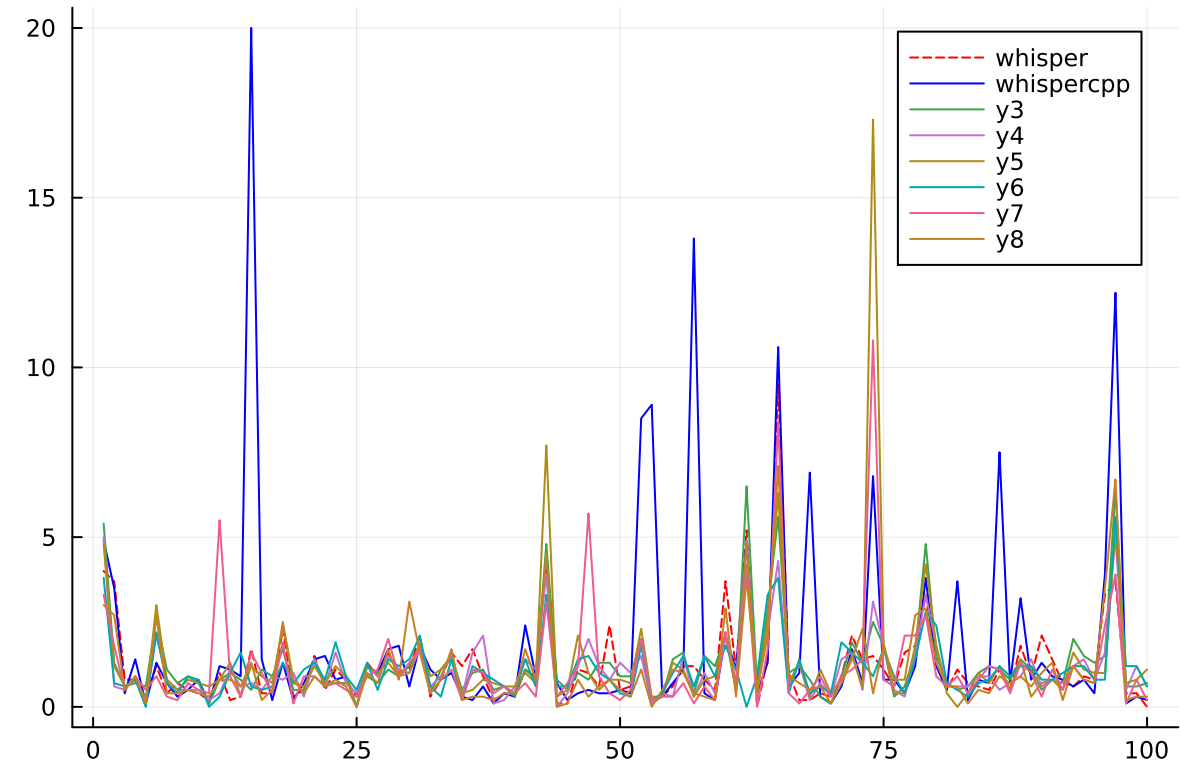
```
1 md"so IBM seems to have an issue with deletions, maybe due to speakers. sorting by
number of speakers shows some effect. The first 68 have a single speaker.
*corrigan* has two speakers, so this doesn't seem totally explanatory for it's high
DER."
```



```
1 begin
2   wers_dn = sort(wers_d, [:nsp])
3   plot(wers_dn.ibm, label="wer")
4   # plot!(wers_d.ibm)
5   plot!(wers_dn.nsp, label="nsp")
6   # plot!(wers.snr, label="snr")
7 end
```

Whisper and Whisper cpp specifically

whispercpp (dark blue solid) stands out among IERs due to hallucinations, but *whisper* (red dashed) does not. *whispercpp* doesn't have the new options to limit hallucinations



```
1 begin
2     wm = wers_i
3     plot(wm.whisper, label="whisper", color=:red, linestyle=:dash)
4     plot!(wm.whispercpp, label="whispercpp", color=:blue)
5     plot!(wm.aws)
6     plot!(wm.azure)
7     plot!(wm.google)
8     plot!(wm.ibm)
9     plot!(wm.nemo)
10    plot!(wm.rev)
11 end
```

Google specifically

```
1 md"### Google specifically"
```

phillytalk10 is an outlier for google with a very high IER. There's a long string of individual digits inserted during a period of silence (no speech), similar to a hallucination in whisper, although it's only digits.

	file	aws	azure	google	ibm	nemo	rev	whisper	whispercpp
1	"phillytalks10"	2.5	3.1	17.3	0.9	10.8	0.4	1.5	6.8

```
1 wers_i[wers_i.file .== "phillytalks10", :]
```

